

task2

May 15, 2024

1 Task 2

A Convolutional Neural Network (CNN) is a specialized type of deep neural network designed for processing grid-like data structures, most commonly images. A CNN consists of several key components: an input layer that receives the raw data, convolutional layers that apply filters to detect local patterns, activation functions like ReLU to introduce non-linearity, and pooling layers to down-sample the feature maps, reducing their dimensionality while retaining important information. After several convolutional and pooling operations, fully connected layers are used to perform high-level reasoning, culminating in an output layer that produces the final predictions, often using a softmax activation for classification tasks. CNNs leverage local receptive fields and parameter sharing, which make them highly effective at learning hierarchical features. These networks have revolutionized the field of computer vision and are widely applied in tasks such as image recognition, object detection, and, increasingly, in cybersecurity for applications like network intrusion detection.

```
[22]: import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

# For demonstration, we create a random dataset
np.random.seed(42)
data_size = 1000
feature_size = 20
data = pd.DataFrame(np.random.rand(data_size, feature_size),
    columns=[f'feature_{i}' for i in range(feature_size)])
data['label'] = np.random.randint(0, 2, size=data_size)

# Output 10 rows from the generated data
print("Sample data (10 rows):")
print(data.head(10))

# Reshape data for CNN input (e.g., treat features as 1D spatial data)
X = data.drop('label', axis=1).values.reshape(-1, feature_size, 1)
y = data['label']
```

```

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)

# Standardize the data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train.reshape(-1, feature_size)).reshape(-1,
↳feature_size, 1)
X_test = scaler.transform(X_test.reshape(-1, feature_size)).reshape(-1,
↳feature_size, 1)

# Define the CNN model
model = tf.keras.Sequential([
    tf.keras.layers.Conv1D(64, kernel_size=3, activation='relu',
↳input_shape=(feature_size, 1)),
    tf.keras.layers.MaxPooling1D(pool_size=2),
    tf.keras.layers.Conv1D(128, kernel_size=3, activation='relu'),
    tf.keras.layers.MaxPooling1D(pool_size=2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.5), # Adding dropout for regularization
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
↳metrics=['accuracy'])

# Train the model
history = model.fit(X_train, y_train, epochs=50, batch_size=32,
↳validation_split=0.2)

# Evaluate the model
y_pred = (model.predict(X_test) > 0.5).astype(int)
accuracy = accuracy_score(y_test, y_pred)
print(f'CNN Accuracy: {accuracy * 100:.2f}%')

```

Sample data (10 rows):

	feature_0	feature_1	feature_2	feature_3	feature_4	feature_5	\
0	0.374540	0.950714	0.731994	0.598658	0.156019	0.155995	
1	0.611853	0.139494	0.292145	0.366362	0.456070	0.785176	
2	0.122038	0.495177	0.034389	0.909320	0.258780	0.662522	
3	0.388677	0.271349	0.828738	0.356753	0.280935	0.542696	
4	0.863103	0.623298	0.330898	0.063558	0.310982	0.325183	
5	0.031429	0.636410	0.314356	0.508571	0.907566	0.249292	
6	0.807440	0.896091	0.318003	0.110052	0.227935	0.427108	
7	0.962447	0.251782	0.497249	0.300878	0.284840	0.036887	

8	0.367783	0.632306	0.633530	0.535775	0.090290	0.835302
9	0.341066	0.113474	0.924694	0.877339	0.257942	0.659984

	feature_6	feature_7	feature_8	feature_9	...	feature_11	feature_12	\
0	0.058084	0.866176	0.601115	0.708073	...	0.969910	0.832443	
1	0.199674	0.514234	0.592415	0.046450	...	0.170524	0.065052	
2	0.311711	0.520068	0.546710	0.184854	...	0.775133	0.939499	
3	0.140924	0.802197	0.074551	0.986887	...	0.198716	0.005522	
4	0.729606	0.637557	0.887213	0.472215	...	0.713245	0.760785	
5	0.410383	0.755551	0.228798	0.076980	...	0.161221	0.929698	
6	0.818015	0.860731	0.006952	0.510747	...	0.222108	0.119865	
7	0.609564	0.502679	0.051479	0.278646	...	0.239562	0.144895	
8	0.320780	0.186519	0.040775	0.590893	...	0.016588	0.512093	
9	0.817222	0.555201	0.529651	0.241852	...	0.897216	0.900418	

	feature_13	feature_14	feature_15	feature_16	feature_17	feature_18	\
0	0.212339	0.181825	0.183405	0.304242	0.524756	0.431945	
1	0.948886	0.965632	0.808397	0.304614	0.097672	0.684233	
2	0.894827	0.597900	0.921874	0.088493	0.195983	0.045227	
3	0.815461	0.706857	0.729007	0.771270	0.074045	0.358466	
4	0.561277	0.770967	0.493796	0.522733	0.427541	0.025419	
5	0.808120	0.633404	0.871461	0.803672	0.186570	0.892559	
6	0.337615	0.942910	0.323203	0.518791	0.703019	0.363630	
7	0.489453	0.985650	0.242055	0.672136	0.761620	0.237638	
8	0.226496	0.645173	0.174366	0.690938	0.386735	0.936730	
9	0.633101	0.339030	0.349210	0.725956	0.897110	0.887086	

	feature_19	label
0	0.291229	1
1	0.440152	1
2	0.325330	1
3	0.115869	0
4	0.107891	0
5	0.539342	0
6	0.971782	1
7	0.728216	0
8	0.137521	1
9	0.779876	0

[10 rows x 21 columns]

Epoch 1/50

```
/opt/anaconda3/lib/python3.11/site-
packages/keras/src/layers/convolutional/base_conv.py:99: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
super().__init__()
```

18/18 0s 4ms/step -
 accuracy: 0.5397 - loss: 0.7015 - val_accuracy: 0.5000 - val_loss: 0.6912
 Epoch 2/50
 18/18 0s 1ms/step -
 accuracy: 0.5376 - loss: 0.6896 - val_accuracy: 0.5143 - val_loss: 0.6937
 Epoch 3/50
 18/18 0s 1ms/step -
 accuracy: 0.5478 - loss: 0.6859 - val_accuracy: 0.5857 - val_loss: 0.6872
 Epoch 4/50
 18/18 0s 1ms/step -
 accuracy: 0.5673 - loss: 0.6834 - val_accuracy: 0.5000 - val_loss: 0.6911
 Epoch 5/50
 18/18 0s 1ms/step -
 accuracy: 0.5908 - loss: 0.6769 - val_accuracy: 0.5143 - val_loss: 0.6866
 Epoch 6/50
 18/18 0s 2ms/step -
 accuracy: 0.6049 - loss: 0.6740 - val_accuracy: 0.5071 - val_loss: 0.6918
 Epoch 7/50
 18/18 0s 2ms/step -
 accuracy: 0.6758 - loss: 0.6415 - val_accuracy: 0.5714 - val_loss: 0.6862
 Epoch 8/50
 18/18 0s 2ms/step -
 accuracy: 0.6460 - loss: 0.6610 - val_accuracy: 0.4786 - val_loss: 0.6901
 Epoch 9/50
 18/18 0s 2ms/step -
 accuracy: 0.6570 - loss: 0.6457 - val_accuracy: 0.5500 - val_loss: 0.6822
 Epoch 10/50
 18/18 0s 2ms/step -
 accuracy: 0.6680 - loss: 0.6225 - val_accuracy: 0.5214 - val_loss: 0.6828
 Epoch 11/50
 18/18 0s 2ms/step -
 accuracy: 0.6970 - loss: 0.6161 - val_accuracy: 0.5286 - val_loss: 0.6812
 Epoch 12/50
 18/18 0s 2ms/step -
 accuracy: 0.6983 - loss: 0.6092 - val_accuracy: 0.5143 - val_loss: 0.6870
 Epoch 13/50
 18/18 0s 2ms/step -
 accuracy: 0.6868 - loss: 0.6083 - val_accuracy: 0.5429 - val_loss: 0.6850
 Epoch 14/50
 18/18 0s 2ms/step -
 accuracy: 0.7549 - loss: 0.5723 - val_accuracy: 0.5571 - val_loss: 0.6927
 Epoch 15/50
 18/18 0s 2ms/step -
 accuracy: 0.7857 - loss: 0.5232 - val_accuracy: 0.5357 - val_loss: 0.6948
 Epoch 16/50
 18/18 0s 2ms/step -
 accuracy: 0.7662 - loss: 0.5215 - val_accuracy: 0.5429 - val_loss: 0.6849
 Epoch 17/50

18/18 0s 2ms/step -
 accuracy: 0.7505 - loss: 0.5093 - val_accuracy: 0.5786 - val_loss: 0.7483
 Epoch 18/50
 18/18 0s 2ms/step -
 accuracy: 0.7381 - loss: 0.5152 - val_accuracy: 0.5643 - val_loss: 0.7401
 Epoch 19/50
 18/18 0s 2ms/step -
 accuracy: 0.7753 - loss: 0.4731 - val_accuracy: 0.5571 - val_loss: 0.7514
 Epoch 20/50
 18/18 0s 2ms/step -
 accuracy: 0.8493 - loss: 0.4204 - val_accuracy: 0.5429 - val_loss: 0.7458
 Epoch 21/50
 18/18 0s 2ms/step -
 accuracy: 0.8133 - loss: 0.4110 - val_accuracy: 0.5571 - val_loss: 0.7762
 Epoch 22/50
 18/18 0s 2ms/step -
 accuracy: 0.8537 - loss: 0.3983 - val_accuracy: 0.5786 - val_loss: 0.7576
 Epoch 23/50
 18/18 0s 2ms/step -
 accuracy: 0.8727 - loss: 0.3469 - val_accuracy: 0.5786 - val_loss: 0.7973
 Epoch 24/50
 18/18 0s 2ms/step -
 accuracy: 0.8843 - loss: 0.3217 - val_accuracy: 0.5929 - val_loss: 0.7690
 Epoch 25/50
 18/18 0s 2ms/step -
 accuracy: 0.8592 - loss: 0.3199 - val_accuracy: 0.5500 - val_loss: 0.8273
 Epoch 26/50
 18/18 0s 2ms/step -
 accuracy: 0.8873 - loss: 0.2900 - val_accuracy: 0.5786 - val_loss: 0.8609
 Epoch 27/50
 18/18 0s 2ms/step -
 accuracy: 0.9179 - loss: 0.2637 - val_accuracy: 0.5929 - val_loss: 0.8662
 Epoch 28/50
 18/18 0s 1ms/step -
 accuracy: 0.8999 - loss: 0.2575 - val_accuracy: 0.6143 - val_loss: 0.8987
 Epoch 29/50
 18/18 0s 1ms/step -
 accuracy: 0.9468 - loss: 0.2157 - val_accuracy: 0.5929 - val_loss: 0.9271
 Epoch 30/50
 18/18 0s 1ms/step -
 accuracy: 0.9173 - loss: 0.2261 - val_accuracy: 0.5643 - val_loss: 0.9749
 Epoch 31/50
 18/18 0s 2ms/step -
 accuracy: 0.9547 - loss: 0.1907 - val_accuracy: 0.5714 - val_loss: 1.0233
 Epoch 32/50
 18/18 0s 1ms/step -
 accuracy: 0.9731 - loss: 0.1521 - val_accuracy: 0.5857 - val_loss: 1.0271
 Epoch 33/50

18/18 0s 1ms/step -
 accuracy: 0.9705 - loss: 0.1344 - val_accuracy: 0.5857 - val_loss: 1.1377
 Epoch 34/50
 18/18 0s 2ms/step -
 accuracy: 0.9586 - loss: 0.1403 - val_accuracy: 0.5643 - val_loss: 1.1056
 Epoch 35/50
 18/18 0s 1ms/step -
 accuracy: 0.9688 - loss: 0.1167 - val_accuracy: 0.5643 - val_loss: 1.1772
 Epoch 36/50
 18/18 0s 2ms/step -
 accuracy: 0.9713 - loss: 0.1096 - val_accuracy: 0.5786 - val_loss: 1.1556
 Epoch 37/50
 18/18 0s 1ms/step -
 accuracy: 0.9825 - loss: 0.0929 - val_accuracy: 0.5857 - val_loss: 1.2567
 Epoch 38/50
 18/18 0s 1ms/step -
 accuracy: 0.9626 - loss: 0.0987 - val_accuracy: 0.6071 - val_loss: 1.2027
 Epoch 39/50
 18/18 0s 1ms/step -
 accuracy: 0.9867 - loss: 0.0718 - val_accuracy: 0.5786 - val_loss: 1.2234
 Epoch 40/50
 18/18 0s 1ms/step -
 accuracy: 0.9968 - loss: 0.0514 - val_accuracy: 0.6000 - val_loss: 1.2829
 Epoch 41/50
 18/18 0s 2ms/step -
 accuracy: 0.9867 - loss: 0.0761 - val_accuracy: 0.5857 - val_loss: 1.3499
 Epoch 42/50
 18/18 0s 2ms/step -
 accuracy: 0.9972 - loss: 0.0467 - val_accuracy: 0.5786 - val_loss: 1.3762
 Epoch 43/50
 18/18 0s 2ms/step -
 accuracy: 0.9838 - loss: 0.0578 - val_accuracy: 0.5786 - val_loss: 1.3635
 Epoch 44/50
 18/18 0s 2ms/step -
 accuracy: 0.9797 - loss: 0.0667 - val_accuracy: 0.6071 - val_loss: 1.4756
 Epoch 45/50
 18/18 0s 2ms/step -
 accuracy: 0.9887 - loss: 0.0543 - val_accuracy: 0.6071 - val_loss: 1.4191
 Epoch 46/50
 18/18 0s 2ms/step -
 accuracy: 0.9884 - loss: 0.0554 - val_accuracy: 0.5714 - val_loss: 1.4520
 Epoch 47/50
 18/18 0s 2ms/step -
 accuracy: 0.9842 - loss: 0.0554 - val_accuracy: 0.5357 - val_loss: 1.5236
 Epoch 48/50
 18/18 0s 2ms/step -
 accuracy: 0.9925 - loss: 0.0406 - val_accuracy: 0.5643 - val_loss: 1.5168
 Epoch 49/50

```
18/18          0s 2ms/step -  
accuracy: 1.0000 - loss: 0.0235 - val_accuracy: 0.5714 - val_loss: 1.5277  
Epoch 50/50  
18/18          0s 2ms/step -  
accuracy: 0.9995 - loss: 0.0223 - val_accuracy: 0.5857 - val_loss: 1.5626  
10/10          0s 2ms/step  
CNN Accuracy: 51.33%
```

Please use the “Kernel>Restart & Run All” command in Jupyter Notebook and check your results before submitting your homework. Note that I rerun all boxes on my side before grading.